

Ensemble Learning

Aldo Solari



Ensembles

- Ensemble methods are techniques that create multiple models and then combine them to produce improved results
- These models, when used as inputs of ensemble methods, are called “base learners” or “weak learners”
- Ensemble learning is appealing because that it is able to boost weak learners which are slightly better than random guess to strong learners which can make very accurate predictions
- However, ensemble learning increases computation time and reduces interpretability

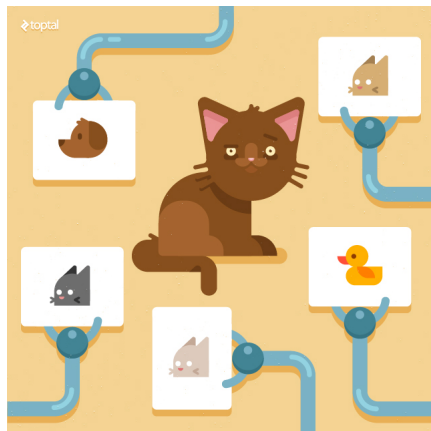


Majority voting

- Majority voting is a simple ensemble method used for classification
- Each classification model in the ensemble makes a prediction (vote) for each test observation and the final prediction is the one that receives more than half of the votes
- If none of the predictions get more than half of the votes, we may say that the ensemble method could not make a stable prediction for this observation



Majority voting



Majority voting

- Suppose that we have three independent classifiers
- Assume that each classifier has probability π of making the correct classification
- The number of correct classifications is

$$C \sim \text{Binomial}(3, \pi)$$

- The majority vote classifier makes the correct prediction when $C \geq 2$, that is, with probability

$$\Pr(C \geq 2) = 3\pi^2(1 - \pi) + \pi^3$$

- e.g. if $\pi = 70\%$, then the majority vote classifier makes the correct prediction with probability 78.4%



Ensemble of trees

- Classification and regression trees are simple and useful for interpretation
- However they are typically not competitive with other approaches in terms of prediction accuracy
- Ensemble methods such as [bagging](#), [random forests](#) and [boosting](#) grow multiple trees which are then combined to yield a single prediction
- Combining a large number of trees can often result in dramatic improvements in prediction accuracy, at the expense of some interpretation loss



Instability of trees

- The primary disadvantage of trees is that they are rather unstable (high variance)
- In other words, a small change in the data often results in a completely different tree
- One major reason for this instability is that if a split changes, all the splits under it change as well, thereby propagating the variability
- Leo Breiman's idea: use the instability!



Outline

① Bagging

② Random Forests

③ Boosting



Bagging

- **Bootstrap aggregation**, or bagging, is a general procedure for reducing the variance of a model and it is particularly useful in the case of regression or classification trees
- Recall that given a set of independent variables Z_1, \dots, Z_B , each with variance σ^2 , the variance of the average

$$\bar{Z} = \frac{1}{B} \sum_{b=1}^B Z_b$$

is σ^2/B

- In other words, averaging a set of variables reduces variance. Of course, this is not practical because we generally do not have multiple training sets
- Instead, we can **bootstrap**, by taking repeated samples from the (single) training set



Bagging algorithm

- First, generate B different bootstrapped training sets

$$(x_1^{*b}, y_1^{*b}), (x_2^{*b}, y_2^{*b}), \dots, (x_n^{*b}, y_n^{*b}) \quad b = 1, \dots, B$$

by sampling with replacement from the training set

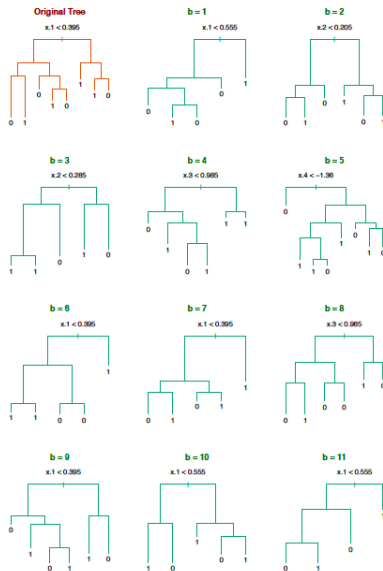
- Then, fit a tree on the b th bootstrapped training set in order to get $\hat{f}^{*b}(x)$, the prediction at a point x
- Finally, average all the predictions to obtain

$$\hat{f}_{bag}(x) = \frac{1}{B} \sum_{b=1}^B \hat{f}^{*b}(x)$$

- For classification trees, we can record the class $\hat{C}^{*b}(x)$ predicted by each of the B trees, and take a majority vote



Many trees



Source: Hastie et al. (2009) p. 284

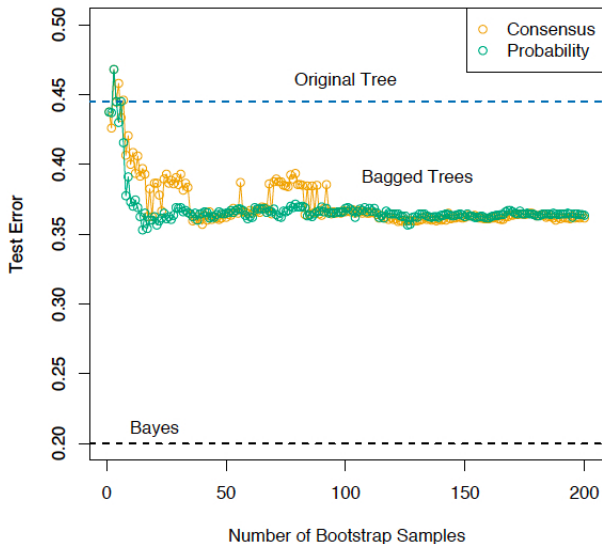


spam data

- 4601 email messages sent to “George” at HP-Labs
- The goal is to build a customized spam filter for George: predict whether an e-mail message is spam (junk mail) or good
- For this problem not all errors are equal; we want to avoid filtering out good e-mail, while letting spam get through is not desirable but less serious in its consequences
- Recorded for each email message is the relative frequency of certain key words (e.g. business, address, free, George) and certain characters: (, [, !, \$, #. Included as well are three different recordings of capitalized letters.



Bagged trees



Source: Hastie et al. (2009) p. 285



Out-of-bag observations

- In the b th bootstrapped training set

$$(x_1^{*b}, y_1^{*b}), (x_2^{*b}, y_2^{*b}), \dots, (x_n^{*b}, y_n^{*b})$$

the probability for one observation not to be drawn in any one of the n draws can be approximated by

$$\lim_{n \rightarrow \infty} \left(1 - \frac{1}{n}\right)^n = \frac{1}{e} \approx 0.368$$

- $\approx 1/3$ of the n original observations are **out-of-bag** (OOB)

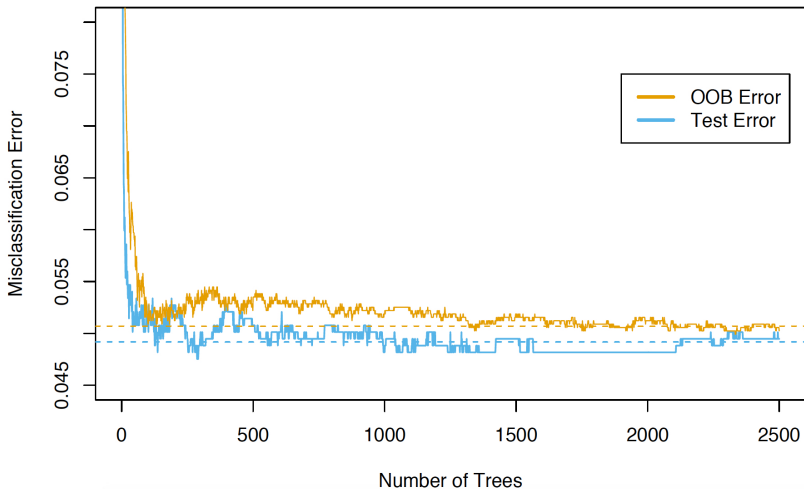


Out-of-bag error

- Each bagged tree makes use of $\approx 2/3$ of the original observations
- We can predict the response for the i th observation using each of the bagged trees in which that observation was OOB
- This yields $\approx B/3$ predictions for the i th observation, which we average/majority vote
- This estimate is essentially the LOOCV error for bagging, if B is large



spam data: OOB error



Source: Hastie et al. (2009) p. 592



Outline

① Bagging

② Random Forests

③ Boosting



Random forests

- Create even more variation in individual trees
- Bagging varies the **rows** of the training set (randomly draw observations)
- Random forests varies also the **columns** of the training set (randomly draw predictors)



Random forests algorithm

- Before each split, select $m \leq p$ of the predictors at random as candidates for splitting
- $m = p$ gives Bagging as a special case
- m is called a **tuning parameter**. Typically $m \approx \sqrt{p}$



De-correlating trees

- Random sampling of the predictors **decorrelates** the trees. This reduces the variance when we average the trees
- Recall that given a set of identical distributed (but not necessarily independent) variables Z_1, \dots, Z_B with pairwise correlation $\mathbb{C}\text{orr}(Z_j, Z_l) = \rho$, mean $\mathbb{E}(Z_j) = \mu$ and variance $\mathbb{V}\text{ar}(Z_j) = \sigma^2$, then

$$\mathbb{V}\text{ar}(\bar{Z}) = \rho\sigma^2 + \frac{(1 - \rho)}{B}\sigma^2 \quad (1)$$

- The idea in random forests is to improve the variance reduction of bagging by reducing the correlation ρ between the trees, without increasing the variance σ^2 too much



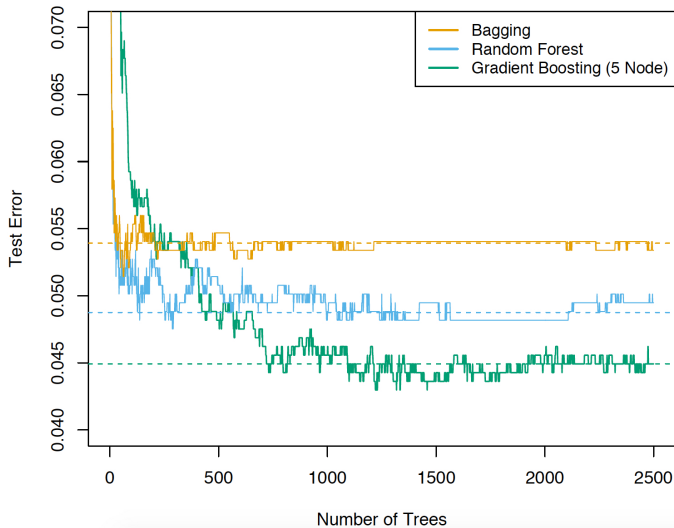
Proof of (1)

- $\rho = \frac{1}{\sigma^2} [\mathbb{E}(Z_i Z_j) - \mathbb{E}(Z_i) \mathbb{E}(Z_j)]$
- $\mathbb{E}(Z_i Z_j) = \rho \sigma^2 + \mu^2$ if $i \neq j$
- $\mathbb{E}(Z_i^2) = \sigma^2 + \mu^2$
- $\mathbb{E}[(\sum_{j=1}^B Z_j)^2] = \sum_{i=1}^B \sum_{j=1}^B \mathbb{E}(Z_i Z_j) = B \mathbb{E}(Z_i^2) + (B^2 - B) \mathbb{E}(Z_i Z_j)$
- $\mathbb{E}(\sum_{j=1}^B Z_j) = \sum_{j=1}^B \mathbb{E}(Z_j) = B\mu$

$$\begin{aligned}\text{Var}(\bar{Z}) &= \frac{1}{B^2} \text{Var}(\sum_{j=1}^B Z_j) \\ &= \frac{1}{B^2} \{ \mathbb{E}[(\sum_{j=1}^B Z_j)^2] - [\mathbb{E}(\sum_{j=1}^B Z_j)]^2 \}\end{aligned}$$



spam data: random forest



Source: Hastie et al. (2009) p. 589



Variable importance

- We can calculate the importance of a predictor X_j
- For each tree, record the accuracy on the OOB observations
- Do the same is done but with X_j values randomly permuted in the OOB observations
- Compute the decrease in each tree's accuracy
- If the average decrease over all the trees is large, then the predictor is considered important - its value makes a big difference in predicting the response
- If the average decrease is small, then the predictor doesn't make much difference to the response

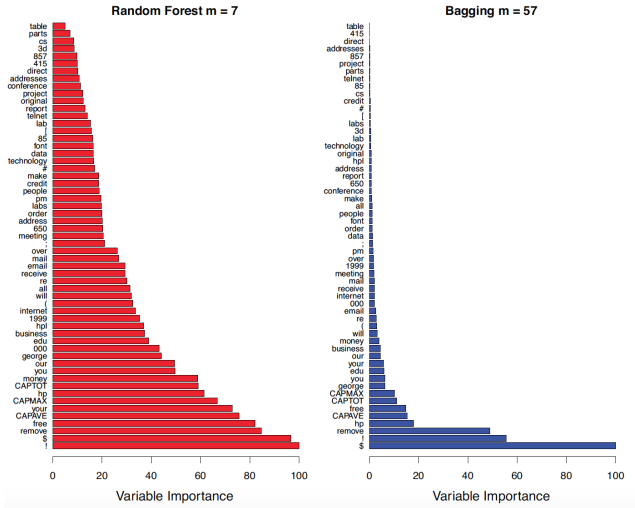


randomForest()

```
fit <- randomForest(y ~ .,  
                    data = train,  
                    ntree = B,  
                    mtry = m,  
                    importance=TRUE)
```



spam data: variable importance



Source: Efron and Hastie (2016) p. 332



Bagging and random forest takeaways

- Bagging stabilizes decision trees and improves accuracy by reducing variance
- Random forests further improve decision tree performance by de-correlating the individual trees in the bagging ensemble
- Random forests' variable importance measures can help you determine which variables are contributing the most strongly to your model
- Because the trees in a random forest ensemble are unpruned and potentially quite deep, there's still a danger of overfitting



Outline

① Bagging

② Random Forests

③ Boosting



Boosting

- Boosting starts by fitting a base learner to the training data
- Next the base learner is re-fitted, but with more weight (importance) given to badly fitted/misclassified observations
- This process is repeated until some stopping rule is reached
- We discuss the following two algorithms
 - boosting with squared error loss (L2-boosting) by using regression trees as weak learners
 - boosting with exponential loss (AdaBoost.M1) by using classification trees as weak learners
- For a more general formulation (gradient boosting), see EH, p. 341 Algorithm 17.4



L2-boosting algorithm for regression trees

1. Initialize $\hat{f}(x) = \bar{y}$ and $r_i = y_i - \bar{y}$ for $i = 1, \dots, n$
2. For $b = 1, 2, \dots, B$, repeat:
 - (a) Fit a tree \hat{f}^b with d splits to the data $(x_1, r_1), \dots, (x_n, r_n)$
 - (b) Update \hat{f} by adding in a shrunk version of the new tree:

$$\hat{f}(x) \leftarrow \hat{f}(x) + \lambda \hat{f}^b(x)$$


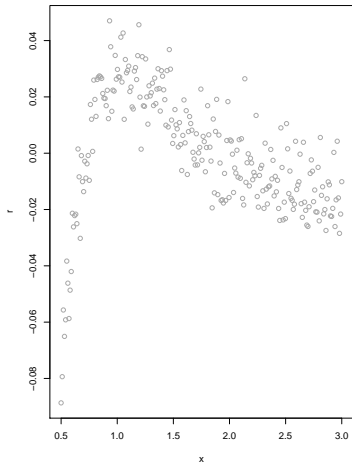
- (c) Update the residuals:



$$r_i \leftarrow r_i - \lambda \hat{f}^b(x_i)$$


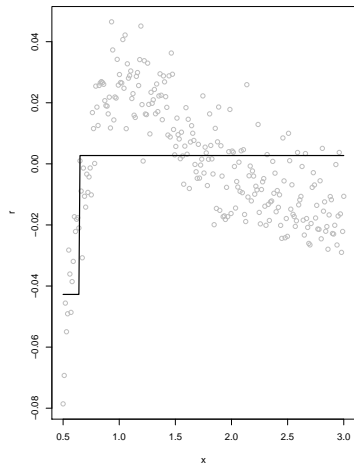
3. Output the boosted model:


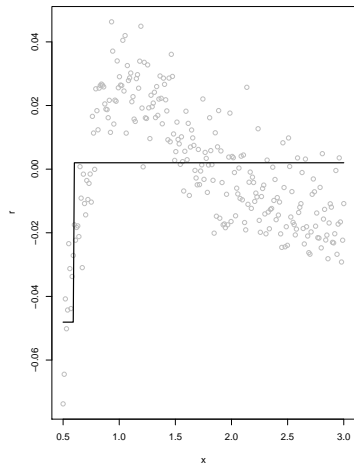
$$\hat{f}(x) = \bar{y} + \sum_{b=1}^B \lambda \hat{f}^b(x)$$


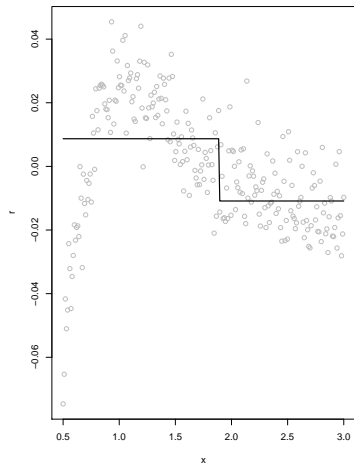



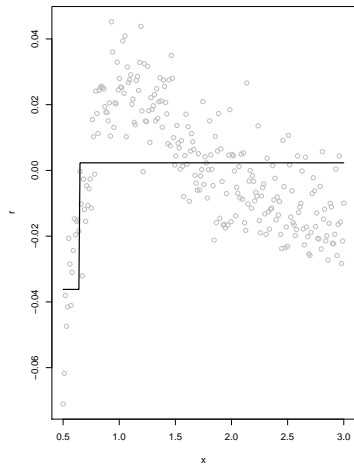



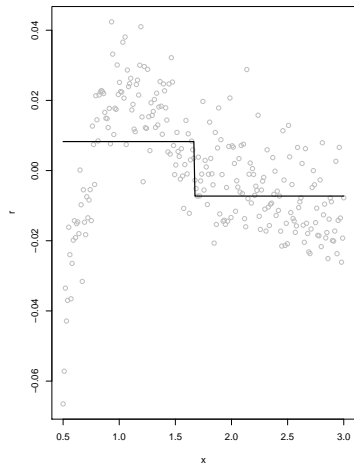


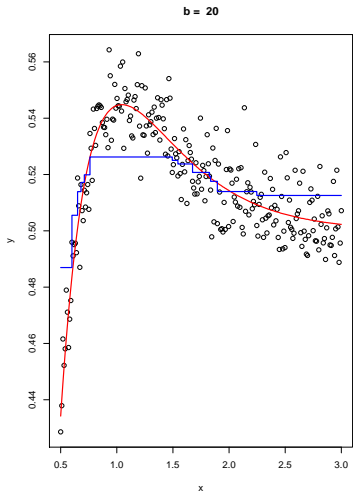
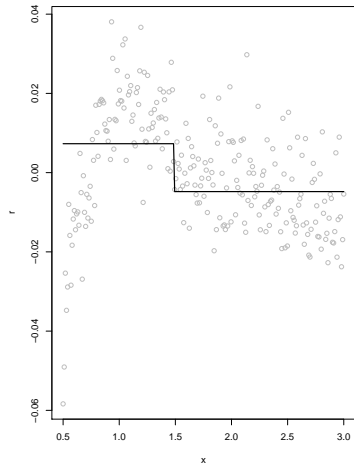






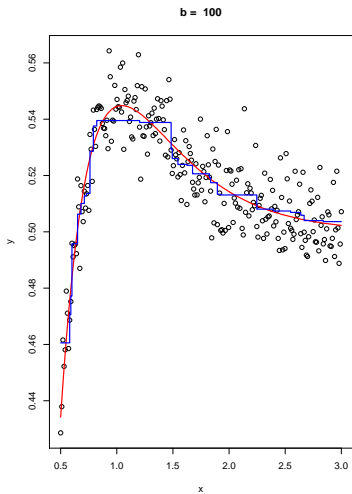
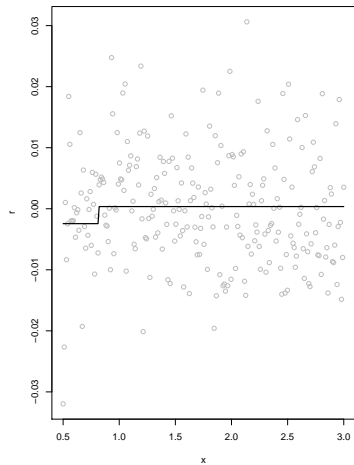
















Tuning parameters for boosting

- The **number of trees** B

Unlike bagging and random forests, boosting can overfit if B is too large, although this overfitting tends to occur slowly if at all. Use cross-validation to select B

- The **shrinkage parameter** λ

A small positive number, which controls the rate at which boosting learns. Typical values are 0.01 or 0.001, and the right choice can depend on the problem. Very small λ can require using a very large value of B in order to achieve good performance

- The **number of splits** d

It controls the complexity of the boosted ensemble. Often $d = 1$ works well, in which case each tree is a stump, consisting of a single split and resulting in an additive model. More generally d is the interaction depth, and controls the interaction order of the boosted model, since d splits can involve at most d variables



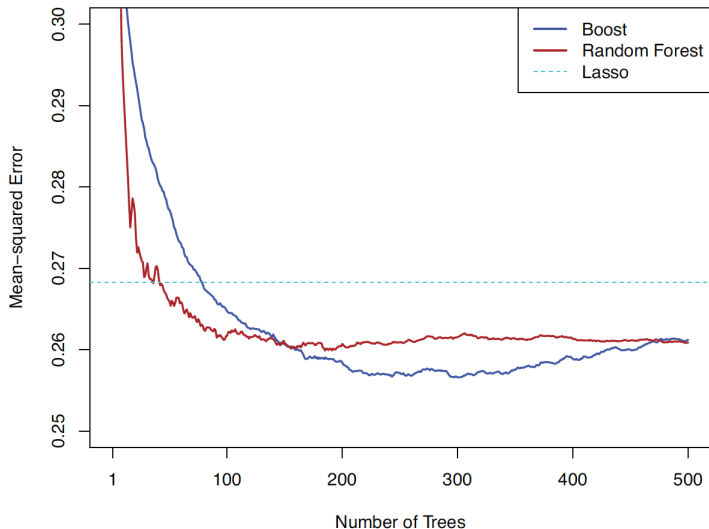
gbm()

```
fit <- gbm(y ~ .,  
           distribution = "gaussian",  
           data = train,  
           n.trees = B,  
           interaction.depth = d,  
           shrinkage = lambda,  
           bag.fraction = 0.5, # default  
           cv.folds = 0) # default
```

- `bag.fraction = 0.5` : grows each new tree on a 50% random sub-sample of the training data. Apart from speeding up the computations, this has a similar effect to bagging, and results in some variance reduction in the ensemble
- `cv.folds = 0` : no cross-validation



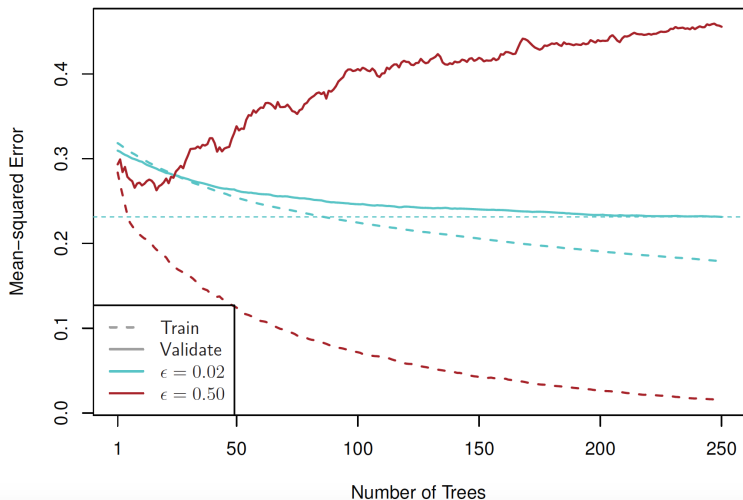
Number of trees B



Source: Efron and Hastie (2016), Figure 17.6



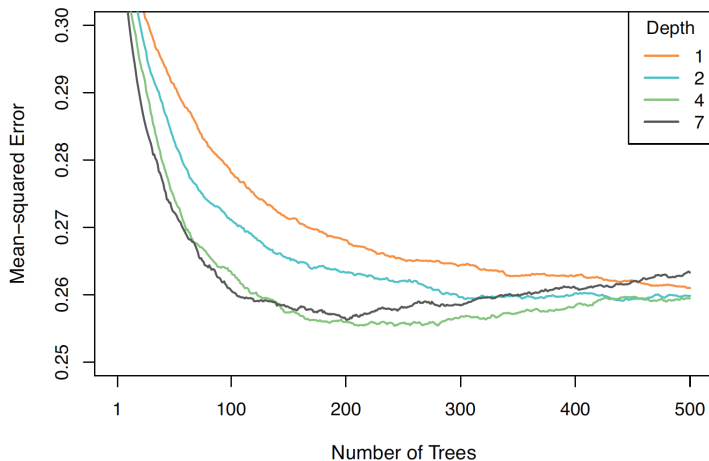
Shrinkage parameter λ



Source: Efron and Hastie (2016), Figure 17.10; ϵ in the legend is λ



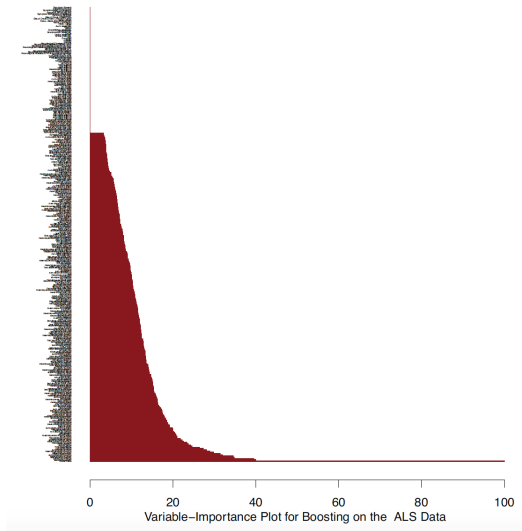
Number of splits d



Source: Efron and Hastie (2016), Figure 17.8



Variable Importance



Source: Efron and Hastie (2016), Figure 17.7



Adaboost: the original boosting algorithm

- Boosting was originally proposed as a means for improving the performance of weak learners in binary classification problems
- The idea is to fit a sequence of classifiers to modified versions of the training data, where the modifications give more weight to misclassified observations
- The final classification is by weighted majority vote
- We present the AdaBoost.M1 algorithm (Freund and Schapire, 1997), which was developed for the response variable coded as $Y \in \{-1, 1\}$ and uses decision trees as weak learners



Adaboost algorithm for classification trees

- ❶ Initialize the observation weights $w_i = 1/n, i = 1, \dots, n$
- ❷ For $b = 1, \dots, B$ repeat the following steps
 - (a) Fit a classification tree \hat{C}^b to the training data, using observation weights w_i
 - (b) Compute the weighted misclassification error for \hat{C}^b

$$\text{Err}^b = \frac{\sum_{i=1}^n w_i I\{y_i \neq \hat{C}^b(x_i)\}}{\sum_{i=1}^n w_i}$$

- (c) Compute

$$\alpha^b = \log[(1 - \text{Err}^b)/\text{Err}^b]$$


- (d) Update the weights as

$$w_i \leftarrow w_i \cdot \exp(\alpha^b \cdot I\{y_i \neq \hat{C}^b(x_i)\}), \quad i = 1, \dots, n$$

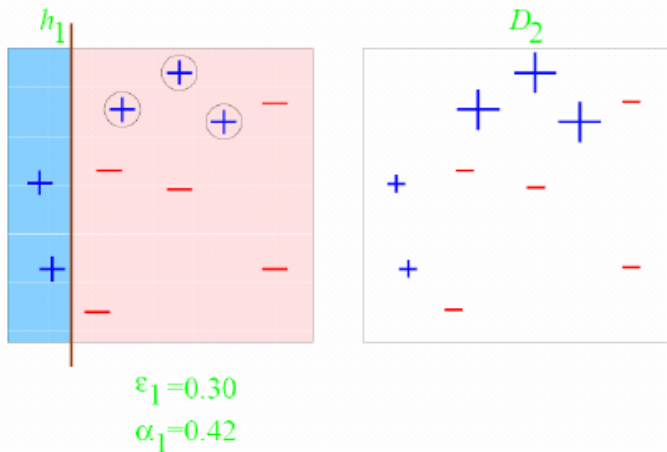
- ❸ Output $\hat{C}^B(x) = \text{sign}(\sum_{b=1}^B \alpha^b \hat{C}^b(x))$



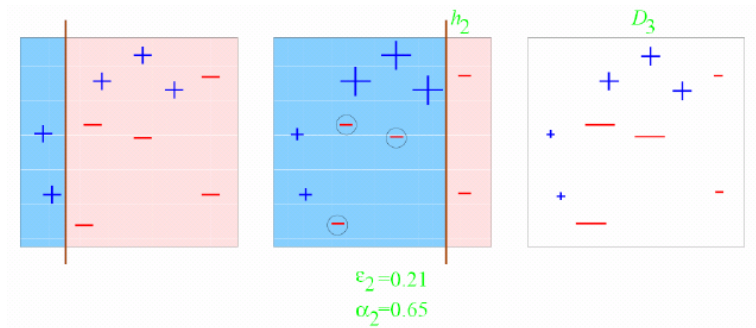
Adaboost example



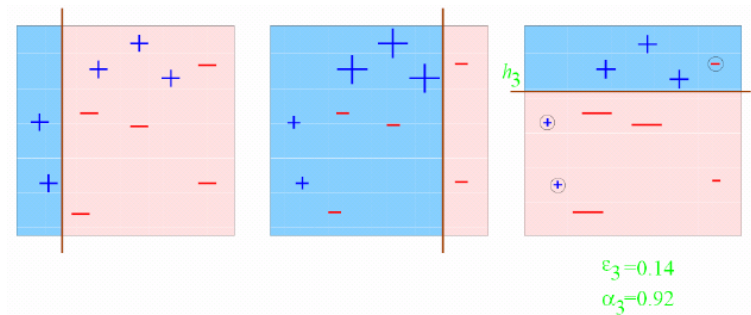
Adaboost example



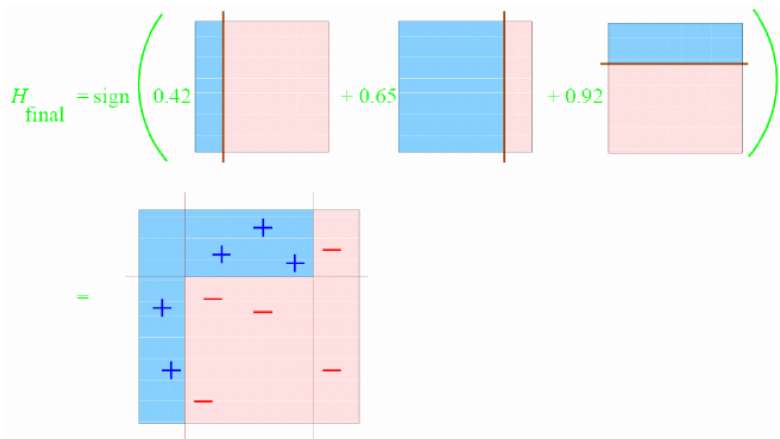
Adaboost example



Adaboost example



Adaboost example



Tuning of boosting

- As is often the case, boosting slightly outperforms a random forest here, but at a price
- Careful tuning of boosting requires considerable extra work, with time-costly rounds of cross-validation, whereas random forests are almost automatic

